

# LVM SNAPSHOT

## What is an LVM Snapshot?

An **LVM snapshot** is a **point-in-time, virtual copy** of a logical volume (LV) created using **copy-on-write (CoW)** technology. It captures the exact state of the volume at the moment of creation and remains consistent even as the original volume continues to be modified.

## How It Works

1. When you create a **snapshot**, it initially shares all data blocks with the origin LV.
2. As the origin is updated, the original data blocks are **copied to the snapshot volume** before the changes occur—ensuring the snapshot reflects the original state.
3. Only changed blocks are stored in the snapshot, making it **space-efficient**
4. New blocks (never before written) aren't copied—only overwritten blocks trigger CoW.

## Why Use Snapshots?

### Consistent Backups with Zero Downtime

Ideal for backing up live systems—snapshots freeze a stable state without disrupting ongoing operations

### Testing & Safe Rollbacks

Perfect for applying patches, upgrades, or experiments on a separate copy. If issues arise, you can rollback effortlessly

## IMPORTANT NOTE:

In LVM we have a feature where an increasing snapshot size **only prevents early fill-up**; it doesn't retroactively include data changed before extension. Only blocks changed *after* creation are tracked and restorable.

## Workflow Example

### Step 1: Check the filesystem(/suresh/data) usage

```
root@localhost:/ # df -hT /suresh/data/
Filesystem      Type  Size  Used Avail Use% Mounted on
/dev/mapper/vg01-lv01  xfs   4.0G  61M  3.9G  2% /suresh/data
root@localhost:/ #
```

This reveals how full the **/suresh/data** filesystem is (e.g., total size, used, available, percentage).

Here, **lv** is **lv01** and **vg** is **vg01**

### Step 2: Add data/files in /suresh/data

Command: **find /etc -name '\*.conf' -exec cp -r {} /suresh/data/ \;**

```
root@localhost:/ # find /etc -name '*conf' -exec cp -r {} /suresh/data/ \;
cp: '/etc/authselect/nsswitch.conf' and '/suresh/data/nsswitch.conf' are the same file
root@localhost:/ #
root@localhost:/ #
root@localhost:/ # ls /suresh/data | wc -l
146
root@localhost:/ # du -sh /suresh/data
1.2M    /suresh/data
root@localhost:/ #
```

## Step 3: Create LV Snapshot

Command:

```
lvcreate -L 1.5G -s -n lv01_snap /dev/vg01/lv01
```

```
lvcreate --size 1.5G --snapshot --name lv01_snap /dev/vg01/lv01
```

```
root@localhost:~# lvcreate --size 100M --snapshot --name lv01_snap /dev/vg01/lv01
Logical volume "lv01_snap" created.
[root@localhost ~]#
[root@localhost ~]# lvdisplay /dev/vg01/lv01
--- Logical volume ---
  LV Path          /dev/vg01/lv01
  LV Name         lv01
  VG Name         vg01
  LV UUID         mbhPqI-5xZz-TcNd-P21q-vYkV-l2vA-3q7dc5
  LV Write Access read/write
  LV Creation host, time localhost.localdomain, 2025-07-21 16:34:02 +0530
  LV snapshot status source of
    lv01_snap [active]
  LV Status        available
  # open           1
  LV Size          4.00 GiB
  Current LE       1024
  Segments         2
  Allocation       inherit
  Read ahead sectors auto
    - currently set to   8192
  Block device     253:2
```

```

root@localhost:~]# lvdisplay /dev/vg01/lv01_snap
--- Logical volume ---
LV Path              /dev/vg01/lv01_snap
LV Name              lv01_snap
VG Name              vg01
LV UUID              CJD1Ww-IEzY-oEsL-yXEk-35LM-eCyh-DRUBxq
LV Write Access      read/write
LV Creation host. time localhost.localdomain. 2025-07-22 10:01:31 +0530
LV snapshot status   active destination for lv01
LV Status            available
# open               0
LV Size              4.00 GiB
Current LE           1024
COW-table size       100.00 MiB
COW-table LE          25
Allocated to snapshot 0.00%
Snapshot chunk size  4.00 KiB
Segments             1
Allocation           inherit
Read ahead sectors   auto
- currently set to   8192
Block device         253:5

```

```

root@localhost:suresh/data]# lvs
  LV      VG Attr  LSize  Pool Origin Data%  Meta%  Move Log Cpy%Sync Convert
  root    rhel -wi-ao--- 16.41g
  swap    rhel -wi-ao--- 2.00g
  lv01    vg01 owi-aos--- 4.00g
  lv01_snap vg01 swi-a-s--- 100.00m    lv01     0.08
[root@localhost data]#
[root@localhost data]# lvs /dev/vg01/lv01_snap
  LV      VG Attr  LSize  Pool Origin Data%  Meta%  Move Log Cpy%Sync Convert
  lv01_snap vg01 swi-a-s--- 100.00m    lv01     0.08
[root@localhost data]#

```

- **lv01\_snap** is a thin snapshot of **lv01**—thin snapshots consume storage only as the original volume changes.

- At **0.08% data usage**, almost no data has changed since snapshot creation: indicating little write activity or early snapshot state.
- The **100 MB LSize** is the allocated snapshot space; given the low Data%, this is sufficient—but once changes exceed ~100 MB, the snapshot may fail due to space exhaustion.

#### Step 4: Remove the data or unfortunately data will deleted in /suresh/data

Command: **rm -rf /suresh/data/\*.conf**

```
root@localhost:/ [root@localhost /]# rm -rf /suresh/data/*.conf [root@localhost /]# ls /suresh/data dconf [root@localhost /]#
```

#### Step 5: Unmount the LV and Start merge with **lvconvert --merge**

Command:

**umount /suresh/data**

**lvconvert --merge /dev/vg01/lv01\_snap**

```
root@localhost:/# umount /suresh/data
[root@localhost /]#
[root@localhost /]# lvconvert --merge /dev/vg01/lv01_snap
Merging of volume vg01/lv01_snap started.
vg01/lv01: Merged: 100.00%
[root@localhost /]#
```

This initiates merging the snapshot (**lv01\_snap**) back into the origin LV (**lv01**).

Unmount first: As done (**umount /suresh/data**), this allows merge to start immediately.

#### How **lvconvert --merge** works

- The command merges the snapshot's data into the origin LV.
- Both the snapshot and the origin LV must be closed/unmounted for immediate merge; otherwise, it's deferred until the next activation.
- After merge:
  - Snapshot LV is deleted.
  - Original LV retains its name, UUID, and metadata

#### Step 6: mount the lv and check the data in /suresh/data

Command:

```
root@localhost:/# mount /dev/vg01/lv01 /suresh/data
[root@localhost /]#
[root@localhost /]# ls /suresh/data/ | wc -l
146
[root@localhost /]#
```

- After merging the snapshot back into the origin (from your previous steps), you remounted **lv01** to verify the current contents.
- Seeing **146 items** confirms that:
  - The snapshot merge was successful.
  - The filesystem is intact and showing expected contents.
- If you know the baseline count from before the snapshot, you can compare to ensure **no data was lost or added unexpectedly**.

```
root@localhost:/#
[root@localhost /]# lvs
  LV   VG Attr       LSize  Pool Origin Data%  Meta%  Move Log Cpy%Sync Convert
  root  rhel -wi-ao---- 16.41g
  swap  rhel -wi-ao----  2.00g
  lv01  vg01 -wi-ao----  4.00g
[root@localhost /]#
[root@localhost /]# lvdisplay /dev/vg01/lv01_snap
  Failed to find logical volume "vg01/lv01_snap"
[root@localhost /]# 
```

## Manual Snapshot Extension

### Check current snapshot usage

Command:

[\*\*lvs -o vg\\_name,lv\\_name,origin,data\\_percent,lv\\_size\*\*](#)

This reports how much of the snapshot's allocated CoW space is used

### Deactivate snapshot (if required)

If the snapshot is active, deactivate it before resizing:

**Command:**

`lvchange -an /dev/vg01/lv01_snap`

### Extend snapshot size

Add more CoW area using:

`lvextend --size +<AdditionalSize>G /dev/<vg>/<snapshot_lv>`

e.g., `lvextend --size +1G /dev/vg01/lv01_snap`

### Verify the extension

**Command:**

`lvs -o lv_name,lv_size,data_percent`

Ensure the snapshot size increased and data% adjusted accordingly

## Conclusion & Key Takeaways

An **LVM snapshot** is a powerful, point-in-time copy of a logical volume leveraging copy-on-write (CoW) technology. It enables consistent backups and risk-free testing environments without disrupting live operations. However, snapshots are **not a replacement for backups**—they require proper space management and monitoring, and are best used as part of a broader backup and recovery strategy.

## Best Practices

- **Create snapshots during quiet times**, ideally after quiescing applications (e.g. databases) to ensure consistency.
- **Monitor snapshot usage (data%)** regularly with `lvs`, and resize before it nears capacity to avoid failure.
- **Use snapshots as backup staging areas**, not long-term solutions. Transfer data off-volume rapidly, then `lvremove` the snapshot.
- **Minimize snapshot count**, as each adds I/O overhead.

### LVM snapshots are excellent tools for:

1. **Zero-downtime backups** — capture consistent volume state.
2. **Safe testing and rollback** — revert to known-good state easily.
3. **Efficient storage** — only track changes, not entire data.

But they must be **monitored for space**, **managed wisely**, and always used in tandem with a robust backup policy.