

RAID:

RAID (Redundant Array of Independent Disks) in Linux is a technology that combines multiple physical drives into a single logical unit.

What is RAID 0 (Striping)?

- **Definition:** Splits data into chunks and distributes them across two or more disks.
- **Purpose:** Improves performance (I/O throughput).
- **Tradeoff:** No redundancy — failure of any disk causes total data loss.

Why RAID 0?

- Splits data into chunks and distributes (“stripes”) them across two or more disks.
- **Boosts performance**—read/write speeds scale with the number of disks.
- **Lacks redundancy:** if any drive fails, you lose the entire array

To configure **RAID 0 (striping)** on Linux:

1. LVM (Logical Volume Manager)
2. MDADM (Multiple Device Administration)

RAID 0 Striping with LVM – Full Workflow

Step 1: Add New Disks & Initialize as PVs

Once you physically attach or virt- add your disks (e.g., `/dev/sdb`, `/dev/sdc` & `/dev/sdd`), format them as LVM PVs:

Command: `pvcreate /dev/sdb /dev/sdc /dev/sdd`

```
root@localhost:~]# lsblk | grep -E '^sdb|sdc|sdd$'
sdb          8:16    0      3G  0 disk
sdc          8:32    0      3G  0 disk
sdd          8:48    0      3G  0 disk
[root@localhost ~]#
[root@localhost ~]# pvcreate /dev/sdb /dev/sdc /dev/sdd
  Physical volume "/dev/sdb" successfully created.
  Physical volume "/dev/sdc" successfully created.
  Physical volume "/dev/sdd" successfully created.
[root@localhost ~]#
```

Creates three PVs ready for LVM usage.

Check with:

```
root@localhost:~]# pvs
PV        VG   Fmt Attr PSize  PFree
/dev/sda3  rhel lvm2 a--  18.41g    0
/dev/sdb       lvm2 ---   3.00g  3.00g
/dev/sdc       lvm2 ---   3.00g  3.00g
/dev/sdd       lvm2 ---   3.00g  3.00g
[root@localhost ~]#
```

Step 2: Create a Volume Group (VG)

Combine the PVs into a volume group:

Command: vgcreate appvg /dev/sdb /dev/sdc /dev/sdd

```
root@localhost:~]# vgcreate appvg /dev/sdb /dev/sdc /dev/sdd
  Volume group "appvg" successfully created
[root@localhost:~]#
[root@localhost:~]# vgs appvg
  VG #PV #LV #SN Attr   VSize   VFree
  appvg   3    0    0 wz--n- <8.99g <8.99g
[root@localhost:~]#
```

Combines the PVs into `appvg`

Step 3: Create a Striped (RAID-0) Logical Volume

To stripe across the 3 PVs:

Command: lvcreate -L 8G -i 3 -I 64K -n lv01 appvg

```
root@localhost:~]# lvcreate -L 8G -i 3 -I 64K -n lv01 appvg
  Rounding size 8.00 GiB (2048 extents) up to stripe boundary size 8.00 GiB (2049 extents).
  Logical volume "lv01" created.
[root@localhost:~]#
```

- `-i 3`: uses 3 stripes (one per PV)
- `-I 64K`: stripe chunk size (adjustable)
- `-L 8G`: size of the LV

Maximum stripes cannot exceed the #PVs in VG

Check with:

```

root@localhost:~#
[root@localhost ~]# lvs -a -o lv_name,vg_name,devices appvg/lv01
  LV   VG   Devices
  lv01 appvg /dev/sdb(0) ,/dev/sdc(0) ,/dev/sdd(0)
[root@localhost ~]#
[root@localhost ~]# lvs -a -o+lv_layout,stripes,devices,stripe_size appvg/lv01
  LV   VG   Attr      LSize Pool Origin Data%  Meta% Move Log Cpy%Sync Convert Layout      #Str Devices
          Stripe
  lv01 appvg -wi-a---- 8.00g                               striped      3 /dev/sdb(0) ,/dev/sdc(0)
(0) ,/dev/sdd(0) 64.00k
[root@localhost ~]#
[root@localhost ~]# lvs --segments -o+segtype,stripes,devices appvg/lv01
  LV   VG   Attr      #Str Type      SSize Type      #Str Devices
  lv01 appvg -wi-a----      3 striped  8.00g striped      3 /dev/sdb(0) ,/dev/sdc(0) ,/dev/sdd(0)
[root@localhost ~]#
[root@localhost ~]#

```

- **SegType:** striped
- **Stripes:** should show 3
- **Devices:** lists PVs used

Step 4: Format & Mount the LV

Format the new LV and mount it:

Commands:

```

mkfs.xfs /dev/appvg/lv01
mkdir /data1
mount /dev/appvg/lv01 /data1

```

```

root@localhost:~#
[root@localhost ~]# mkfs.xfs /dev/appvg/lv01
meta-data=/dev/appvg/lv01      isize=512      agcount=16, agsize=131120 blks
                                sectsz=512    attr=2, projid32bit=1
                                =           crc=1        finobt=1, sparse=1, rmapbt=0
                                =           reflink=1   bigtime=1 inobtcount=1 nrext64=0
data     =           bsize=4096   blocks=2097920, imaxpct=25
                                =           sunit=16    swidth=48 blks
naming   =version 2      bsize=4096   ascii-ci=0, ftype=1
log      =internal log   bsize=4096   blocks=16384, version=2
                                =           sectsz=512  sunit=16 blks, lazy-count=1
realtime =none            extsz=4096   blocks=0, rttextents=0
[root@localhost ~]#

```

```
root@localhost:~# blkid /dev/appvg/lv01
/dev/appvg/lv01: UUID="063f95c9-6e0b-469e-9f2f-9db3894741f5" TYPE="xfs"
[root@localhost ~]#
```

```
root@localhost:~
[root@localhost ~]# mkdir /data1
[root@localhost ~]# mount /dev/appvg/lv01 /data1
[root@localhost ~]# df -h /data1
Filesystem           Size   Used  Avail Use% Mounted on
/dev/mapper/appvg-lv01 8.0G   90M  7.9G   2% /data1
[root@localhost ~]#
```

Step 5: Add entry to /etc/fstab for persistent mount:

Create a stable mount entry so that `/data1` mounts automatically at every boot.

Command: vi /etc/fstab

```
root@localhost:~#
[root@localhost ~]# cat /etc/fstab | grep /data1
UUID=063f95c9-6e0b-469e-9f2f-9db3894741f5      /data1  xfs      defaults      0 0
[root@localhost ~]#
[root@localhost ~]# mount | grep /data1
/dev/mapper/appvg-lv01 on /data1 type xfs (rw,relatime,seclabel,attr2,inode64,logbufs=8,logbsize=64k,sunit=128,swidth=384,noquota)
[root@localhost ~]#
[root@localhost ~]#
```

```
root@localhost:~#
[root@localhost ~]# df -h /data1
Filesystem           Size   Used  Avail Use% Mounted on
/dev/mapper/appvg-lv01 8.0G   90M  7.9G   2% /data1
[root@localhost ~]#
```

```
[root@localhost ~]#
```

```
[root@localhost ~]# lsblk /dev/sdc /dev/sdb /dev/sdd
NAME      MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sdb        8:16   0   3G  0 disk 
└─appvg-lv01 253:2   0   8G  0 lvm   /data1
sdc        8:32   0   3G  0 disk 
└─appvg-lv01 253:2   0   8G  0 lvm   /data1
sdd        8:48   0   3G  0 disk 
└─appvg-lv01 253:2   0   8G  0 lvm   /data1
[root@localhost ~]#
```

Final Takeaway

RAID 0 gives you **speed** and **capacity efficiency**, but **no protection** against disk failure. When combined with a solid backup solution, it creates a balanced system: high performance with robust data safety