

## Title 5: RAID and Advanced Filesystem Features in Linux (mdadm, ext4, and XFS Tuning, Performance Test)

---

### What is RAID?

RAID (Redundant Array of Independent Disks) is a technology that combines multiple physical disks into a single logical unit to improve performance, redundancy, or both.

The idea is simple:

Instead of relying on one disk (which can fail or be slow), we use multiple disks together for speed, reliability, or both.

RAID is handled either by:

**Software RAID** → e.g., mdadm on Linux

**Hardware RAID** → via dedicated RAID controller card

---

### RAID Levels (Types)

There are many RAID levels, but these are the most common:

| RAID Level                               | Description   | Minimum Disks | Storage Efficiency | Fault Tolerance | Performance   |
|--|---|---------------|--------------------|-----------------|---|
| <b>RAID 0<br/>(Striping)</b>             | Data split (striped) across disks for speed. No redundancy. | 2             | 100%               | None            |  High read/write           |
| <b>RAID 1<br/>(Mirroring)</b>            | Data duplicated on each disk.                               | 2             | 50%                | 1 disk          |  Good read, slower write   |
| <b>RAID 5<br/>(Striping with Parity)</b> | Data + parity distributed across disks.                     | 3             | (N-1)/N            | 1 disk          |  Good read, moderate write |
| <b>RAID 6<br/>(Double Parity)</b>        | Like RAID 5 but can survive 2 disk failures.                | 4             | (N-2)/N            | 2 disks         |  Moderate                  |
| <b>RAID 10<br/>(1+0)</b>                 | Combination of mirroring + striping.                        | 4             | 50%                | 1 per mirror    |  Excellent overall         |

---

### Purpose of RAID

RAID is mainly used for:

**Data Redundancy** → Prevent data loss if a disk fails.

**Performance Improvement** → Increase read/write speed.

**Larger Logical Volumes** → Combine disks into one large storage.

**High Availability** → Ensure continuous access even after disk failure.

---

## Pros (Advantages)

| Benefit         | Explanation  |
|-----------------|--|
| Fault Tolerance | Some RAID levels (like RAID 1, 5, 6, 10) protect against disk failure.       |
| Improved Speed  | Striping improves read/write performance (especially RAID 0, 5, 10).         |
| Scalability     | Easy to expand storage by adding more disks.                                 |
| Availability    | Systems remain operational even if one disk fails (depending on RAID level). |

## Don'ts / Limitations

| Caution                    | Why  |
|----------------------------|--|
| Don't use RAID as a backup | RAID protects against hardware failure, not accidental deletions or ransomware.                  |
| RAID rebuilds can be risky | Rebuilding after a disk failure can stress other disks, potentially leading to a second failure. |
| RAID 0 has no redundancy   | One disk fails → total data loss.  |
| Requires identical disks   | Mismatched sizes reduce efficiency.  |
| Software RAID consumes CPU | Minimal, but still a factor compared to hardware RAID.   |

## RAID 5 Example (the one you're using)

Needs 3 or more disks.

Combines striping (speed) + parity (redundancy).

If one disk fails, you can rebuild from parity.

You get N-1 usable capacity.

### Example:

If you have three 512 MB disks →

RAID 5 usable =  $512 \text{ MB} \times (3-1) = \sim 1 \text{ GB usable space}$

| Disk Count | Disk Size (Each) | Total Virtual Storage | Notes                                     |
|------------|------------------|-----------------------|---|
| 3 disks    | 500 MB each      | ~1 GB usable          | Good for basic RAID 5 test                |
| 5 disks    | 1 GB each        | ~4 GB usable          | Better for extended practice & benchmarks |

## Step 1: Creating 3 Disks of 500 MB Each

Before adding 3 disks, the output of **lsblk**, **ls -ld /dev/nvmeOn\***.it helps us to compare easily, after adding disks.

```
[root@indlvrh91 ~]# lsblk -f
NAME      FSTYPE   FSVER     LABEL           UUID
sr0       iso9660
nvme0n1
└─nvme0n1p1    vfat      FAT32
└─nvme0n1p2    xfs
└─nvme0n1p3    LVM2_member LVM2 001
  ├─rhel-root   xfs
  └─rhel-swap   swap      1
nvme0n2        LVM2_member LVM2 001
└─vg_data-MyScripts ext4      1.0
└─vg_data-Mintu   ext4      1.0
nvme0n3
└─nvme0n3p1    LVM2_member LVM2 001
└─vg_data-Mintu ext4      1.0
└─nvme0n3p2    LVM2_member LVM2 001
[root@indlvrh91 ~]# lsblk
NAME      MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sr0       11:0      1   11G  0 rom
nvme0n1  259:0     0   50G  0 disk
└─nvme0n1p1 259:1     0   600M 0 part /boot/efi
└─nvme0n1p2 259:2     0   1G   0 part /boot
└─nvme0n1p3 259:3     0   48.4G 0 part
  ├─rhel-root 253:0     0   46.4G 0 lvm /
  └─rhel-swap 253:1     0   2G   0 lvm [SWAP]
nvme0n2  259:4     0   50G  0 disk
└─vg_data-MyScripts 253:2     0   30G  0 lvm /myscripts
└─vg_data-Mintu   253:3     0   20G  0 lvm /mintu
nvme0n3  259:5     0   20G  0 disk
└─nvme0n3p1 259:6     0   5G   0 part
└─vg_data-Mintu 253:3     0   20G  0 lvm /mintu
└─nvme0n3p2 259:7     0   5G   0 part
[root@indlvrh91 ~]# ls -ld /dev/nvmeOn*
brw-rw----. 1 root disk 259, 0 Oct 14 19:20 /dev/nvme0n1
brw-rw----. 1 root disk 259, 1 Oct 14 19:21 /dev/nvme0n1p1
brw-rw----. 1 root disk 259, 2 Oct 14 19:20 /dev/nvme0n1p2
brw-rw----. 1 root disk 259, 3 Oct 14 19:20 /dev/nvme0n1p3
brw-rw----. 1 root disk 259, 4 Oct 14 19:20 /dev/nvme0n2
brw-rw----. 1 root disk 259, 5 Oct 14 19:20 /dev/nvme0n3
brw-rw----. 1 root disk 259, 6 Oct 14 19:20 /dev/nvme0n3p1
brw-rw----. 1 root disk 259, 7 Oct 14 19:20 /dev/nvme0n3p2
[root@indlvrh91 ~]# date
Tuesday 14 October 2025 07:26:17 PM IST
[root@indlvrh91 ~]#
```

## Verify Disks

First, make sure these disks are **not in use or mounted**.

```
#lsblk; #blkid
```

```
[root@indlvrh91 ~]# lsblk
NAME      MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sr0       11:0      1   11G  0 rom
nvme0n1  259:0     0   50G  0 disk
└─nvme0n1p1 259:1     0   600M 0 part /boot/efi
└─nvme0n1p2 259:2     0   1G   0 part /boot
└─nvme0n1p3 259:3     0   48.4G 0 part
  ├─rhel-root 253:0     0   46.4G 0 lvm /
  └─rhel-swap 253:1     0   2G   0 lvm [SWAP]
nvme0n2  259:4     0   50G  0 disk
nvme0n3  259:5     0   512M 0 disk
nvme0n4  259:6     0   512M 0 disk
nvme0n5  259:7     0   512M 0 disk
[root@indlvrh91 ~]# fdisk -l /dev/nvme0n3 /dev/nvme0n4 /dev/nvme0n5
Disk /dev/nvme0n3: 512 MiB, 536870912 bytes, 1048576 sectors
Disk model: VMware Virtual NVMe Disk
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes

Disk /dev/nvme0n4: 512 MiB, 536870912 bytes, 1048576 sectors
Disk model: VMware Virtual NVMe Disk
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes

Disk /dev/nvme0n5: 512 MiB, 536870912 bytes, 1048576 sectors
Disk model: VMware Virtual NVMe Disk
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
[root@indlvrh91 ~]# date
Tuesday 14 October 2025 07:43:57 PM IST
[root@indlvrh91 ~]#
```

If any of those devices (nvme0n3, nvme0n4, nvme0n5) show partitions or filesystems, **wipe them** **sssclean**:

```
#wipesfs -a /dev/nvme0n3  
#wipesfs -a /dev/nvme0n4  
#wipesfs -a /dev/nvme0n5  
[root@indlvrh91 ~]# wipesfs -a /dev/nvme0n3  
[root@indlvrh91 ~]# wipesfs -a /dev/nvme0n4  
[root@indlvrh91 ~]# wipesfs -a /dev/nvme0n5  
[root@indlvrh91 ~]#
```

Now, three disks /dev/nvme0n3, /dev/nvme0n4, and /dev/nvme0n5 are **completely clean** and ready to be used for your RAID 5 array.

## Step 2: Install Required Packages

**Make sure mdadm (RAID tool) and xfsprogs / e2fsprogs are installed:**

```
#dnf install -y mdadm xfsprogs e2fsprogs
```

Here, I have already installed all the supported packages at the time of patching.

## Using Yum/dnf list installed

```
# dnf list installed mdadm xfsprogs e2fsprogs
```

## **Expected output:**

You should see something like this if installed:

## **Installed Packages**

```
mdadm.x86_64      4.3-4.el9_5 @rhel-9-baseos
xfsprogs.x86_64   6.4.0-5.el9   @rhel-9-baseos
e2fsprogs.x86_64  1.46.5-7.el9  @rhel-9-baseos
```

**Using rpm -q (checks individually)**  
`#rpm -q mdadm`

```
#rpm -q madam  
#rpm -q vfsprogs
```

```
#rpm -q xfsprogs  
#rpm -q xfsprogs
```

```
#rpm -q e2fsprogs
```

## **Expected output if installed:**

- e2fsprogs → version 1.46.5-7.el9
  - mdadm → version 4.3-4.el9\_5
  - xfsprogs → version 6.4.0-5.el9

### Check if the binaries are available

```
#which mdadm; #which mkfs.xfs; #which mkfs.ext4
```

### Expected output:

- /usr/sbin/mdadm
- /usr/sbin/mkfs.xfs
- /usr/sbin/mkfs.ext4

```
[root@indlvrh91 ~]#
[root@indlvrh91 ~]# which mdadm;which mkfs.xfs;which mkfs.ext4
/usr/sbin/mdadm
/usr/sbin/mkfs.xfs
/usr/sbin/mkfs.ext4
[root@indlvrh91 ~]# date
Wednesday 15 October 2025 10:31:59 AM IST
[root@indlvrh91 ~]# █
```

---

### Step 3: Create RAID 5 Array

We'll use /dev/nvme0n3, /dev/nvme0n4, /dev/nvme0n5.

```
#mdadm --create --verbose /dev/md0 --level=5 --raid-devices=3 /dev/nvme0n3 /dev/nvme0n4
/dev/nvme0n5
```

#### ***Detail explanation of command***

| Part   | Meaning / Function   | Explanation  |
|--|----------------------|--|
| mdadm  | RAID management tool | The Linux utility used to create, assemble, monitor, and manage software RAID arrays.            |
| --create                                     | Create new array     | Tells mdadm to create a new RAID device. This erases existing data on listed disks.              |
| --verbose                                    | Detailed output      | Enables verbose mode, showing what mdadm is doing — useful for learning or troubleshooting.      |
| /dev/md0                                     | Target RAID device   | The logical RAID device that will be created. You can name it /dev/md0, /dev/md1, etc.           |
| --level=5                                    | RAID level           | Specifies RAID 5 (striped with parity) — combines speed + redundancy, requires $\geq 3$ disks.   |
| --raid-devices=3                             | Number of disks      | Defines how many physical drives participate in this RAID array. In this case: 3 disks.          |
| /dev/nvme0n3<br>/dev/nvme0n4<br>/dev/nvme0n5 | Physical disks       | Lists the actual disk devices used to build the RAID. These must be wiped clean before creation. |

---

### Quick Calculation

| Total Disks | RAID Level | Formula             | Usable Capacity             | Fault Tolerance                   |
|-------------|------------|---------------------|-----------------------------|-----------------------------------|
| 3 × 512 MB  | RAID 5     | (N - 1) × Disk Size | (3 - 1) × 512 MB = ~1<br>GB | Can survive <b>1 disk failure</b> |

```
[root@indlvrh91:~]# mdadm --create --verbose /dev/md0 --level=5 --raid-devices=3 /dev/nvme0n3 /dev/nvme0n4 /dev/nvme0n5
mdadm: layout defaults to left-symmetric
mdadm: layout defaults to left-symmetric
mdadm: chunk size defaults to 512K
mdadm: size set to 522240K
mdadm: Defaulting to version 1.2 metadata
mdadm: array /dev/md0 started.
[root@indlvrh91 ~]# cat /proc/mdstat
```

### Monitor build progress:

```
#watch cat /proc/mdstat
```

### Optional Verification After Creation

```
#cat /proc/mdstat      # Shows the RAID progress and sync status
#mdadm --detail /dev/md0  # Shows RAID level, devices, state, UUID
#lsblk          # Shows new md device with size
```

```
[root@indlvrh91 ~]# date
Wednesday 15 October 2025 11:33:18 AM IST
[root@indlvrh91 ~]# cat /proc/mdstat
Personalities : [raid6] [raid5] [raid4]
md0 : active raid5 nvme0n3[4] nvme0n5[3] nvme0n4[1]
      1044480 blocks super 1.2 level 5, 512k chunk, algorithm 2 [3/3] [UUU]

unused devices: <none>
[root@indlvrh91 ~]# mdadm --detail /dev/md0
/dev/md0:
      Version : 1.2
      Creation Time : Tue Oct 14 19:54:26 2025
      Raid Level : raid5
      Array Size : 1044480 (1020.00 MiB 1069.55 MB)
      Used Dev Size : 522240 (510.00 MiB 534.77 MB)
      Raid Devices : 3
      Total Devices : 3
      Persistence : Superblock is persistent

      Update Time : Wed Oct 15 11:22:26 2025
      State : clean
      Active Devices : 3
      Working Devices : 3
      Failed Devices : 0
      Spare Devices : 0

      Layout : left-symmetric
      Chunk Size : 512K

Consistency Policy : resync

      Name : indlvrh91.na.xom.com:0 (local to host indlvrh91.na.xom.com)
      UUID : 02008a21:9cee95b9:5562918f:26cebbed
      Events : 84

      Number  Major  Minor  RaidDevice State
        4      259      5        0     active sync   /dev/nvme0n3
        1      259      6        1     active sync   /dev/nvme0n4
        3      259      7        2     active sync   /dev/nvme0n5
```

```
[root@indlvrh91 ~]# lsblk
NAME      MAJ:MIN RM  SIZE RO TYPE  MOUNTPOINTS
sr0        11:0    1   11G  0 rom
nvme0n1   259:0   0   50G  0 disk
└─nvme0n1p1 259:1   0  600M  0 part  /boot/efi
└─nvme0n1p2 259:2   0   1G  0 part  /boot
└─nvme0n1p3 259:3   0 48.4G  0 part
  ├─rhel-root 253:0   0 46.4G  0 lvm   /
  └─rhel-swap 253:1   0   2G  0 lvm   [SWAP]
nvme0n2   259:4   0   50G  0 disk
nvme0n3   259:5   0   512M 0 disk
└─md0      9:0    0 1020M 0 raid5
nvme0n4   259:6   0   512M 0 disk
└─md0      9:0    0 1020M 0 raid5
nvme0n5   259:7   0   512M 0 disk
└─md0      9:0    0 1020M 0 raid5
[root@indlvrh91 ~]#
```

#### **Step 4: Save RAID Configuration (Persistence)**

Make sure RAID auto-assembles after reboot:

```
#mdadm --detail --scan | sudo tee -a /etc/mdadm.conf
#dracut -f
#cat /etc/mdadm.conf
```

```
Last login: Tue Oct 14 19:40:54 2025 from 192.168.71.1
[root@indlvrh91 ~]# mdadm --detail --scan | sudo tee -a /etc/mdadm.conf
ARRAY /dev/md/0 metadata=1.2 UUID=02008a21:9cee95b9:5562918f:26cebbed
[root@indlvrh91 ~]# dracut -f
[root@indlvrh91 ~]# cat /etc/mdadm.conf
ARRAY /dev/md/0 metadata=1.2 UUID=02008a21:9cee95b9:5562918f:26cebbed
[root@indlvrh91 ~]#
[root@indlvrh91 ~]#
```

This saves the array info and rebuilds the initramfs image with RAID support.

#### **Step 5: Create a Filesystem**

**Create a filesystem** (ext4 or XFS) on /dev/md0:

**Ext4:**

```
#mkfs.ext4 /dev/md0
#mkdir -p /mnt/raid-ext4
#mount /dev/md0 /mnt/raid-ext4
#df -hT /mnt/raid-ext4
```

**XFS (RHEL 9 default), I have chosen.**

**Format /dev/md0 with XFS**

```
#mkfs.xfs /dev/md0
#mkfs.xfs → creates an XFS filesystem on the RAID device.
By default, it uses the full array size and the default allocation settings.
```

## Create a mount point

```
#mkdir -p /mnt/raid-xfs
```

This is where your RAID filesystem will be accessible.

## Mount the XFS filesystem

```
#mount /dev/md0 /mnt/raid-xfs
```

The RAID array is now ready for use under /mnt/raid-xfs.

```
root@indlvrh91:~# mkfs.xfs /dev/md0
mkfs.xfs: small data volume, ignoring data volume stripe unit 1024 and stripe width 2048
meta-data=/dev/md0          isize=512    agcount=4, agsize=65280 blks
                          =         sectsz=512   attr=2, projid32bit=1
                          =         crc=1     finobt=1, sparse=1, rmapbt=0
data      =         reflink=1  bigtime=1 inobtcount=1 nrext64=0
          bsize=4096   blocks=261120, imaxpct=25
          sunit=0     swidth=0 blks
naming    =version 2        bsize=4096   ascii-ci=0, ftype=1
log       =internal log     bsize=4096   blocks=16384, version=2
          =         sectsz=512   sunit=0 blks, lazy-count=1
realtime  =none             extsz=4096   blocks=0, rtextents=0
[root@indlvrh91 ~]#
[root@indlvrh91 ~]# mkdir -p /mnt/raid-xfs
[root@indlvrh91 ~]#
[root@indlvrh91 ~]# mount /dev/md0 /mnt/raid-xfs
[root@indlvrh91 ~]#
```

## Verify the mount and filesystem

```
#df -hT /mnt/raid-xfs
```

```
#xfs_info /mnt/raid-xfs
```

#df -hT → shows size, used, and filesystem type.

#lsblk -f → shows RAID device and filesystem info.

#xfs\_info → gives detailed info about XFS layout, block size, and metadata.

```
root@indlvrh91:~# df -hT /mnt/raid-xfs
Filesystem  Type  Size  Used Avail Use% Mounted on
/dev/md0    xfs   956M  39M  918M  5% /mnt/raid-xfs
[root@indlvrh91 ~]# lsblk -f
NAME   FSTYPE  FSVER  LABEL           UUID
sr0    iso9660 Joliet Extension RHEL-9-5-0-BaseOS-x86_64 2024-10-07-15-22-00
nvme0n1
└─nvme0n1p1 vfat    FAT32
└─nvme0n1p2 xfs
└─nvme0n1p3 LVM2_member  LVM2 001
  └─rhel-root xfs
  └─rhel-swap swap    1
nvme0n2    LVM2_member  LVM2 001
nvme0n3    linux_raid_member 1.2  indlvrh91.na.xom.com:0
└─md0      xfs
nvme0n4    linux_raid_member 1.2  indlvrh91.na.xom.com:0
└─md0      xfs
nvme0n5    linux_raid_member 1.2  indlvrh91.na.xom.com:0
└─md0      xfs
[root@indlvrh91 ~]# xfs_info /mnt/raid-xfs
meta-data=/dev/md0          isize=512    agcount=4, agsize=65280 blks
                          =         sectsz=512   attr=2, projid32bit=1
                          =         crc=1     finobt=1, sparse=1, rmapbt=0
data      =         reflink=1  bigtime=1 inobtcount=1 nrext64=0
          bsize=4096   blocks=261120, imaxpct=25
          sunit=0     swidth=0 blks
naming    =version 2        bsize=4096   ascii-ci=0, ftype=1
log       =internal log     bsize=4096   blocks=16384, version=2
          =         sectsz=512   sunit=0 blks, lazy-count=1
realtime  =none             extsz=4096   blocks=0, rtextents=0
[root@indlvrh91 ~]# date
Tuesday 14 October 2024 08:10:53 PM IST
[root@indlvrh91 ~]#
```

## Auto-mount on Boot

Add an entry to /etc/fstab:

```
#blkid /dev/md0
```

Copy the UUID of /dev/md0.

Then edit /etc/fstab:

```
#vi /etc/fstab
```

Add:

```
UUID="de33c6fb-692e-47fd-9f9f-6a629b8e291e" /mnt/raid-xfs xfs defaults 0 0
```

This ensures your RAID XFS mounts automatically after a reboot.

```
[root@indlvrh91 ~]# 
[root@indlvrh91 ~]# blkid /dev/md0
/dev/md0: UUID="de33c6fb-692e-47fd-9f9f-6a629b8e291e" TYPE="xfs"
[root@indlvrh91 ~]# vi /etc/fstab
[root@indlvrh91 ~]# cat /etc/fstab

#
# /etc/fstab
# Created by anaconda on Sun Apr 27 12:50:32 2025
#
# Accessible filesystems, by reference, are maintained under '/dev/disk/'.
# See man pages fstab(5), findfs(8), mount(8) and/or blkid(8) for more info.
#
# After editing this file, run 'systemctl daemon-reload' to update systemd
# units generated from this file.
#
/dev/mapper/rhel-root   /           xfs     defaults        0  0
UUID=e1740ed5-031a-4a33-8944-77e8f5c8ccb6 /boot      xfs     defaults        0  0
UUID=8135-844A          /boot/efi    vfat    umask=0077,shortname=winnt 0  2
/dev/mapper/rhel-swap   none       swap    defaults        0  0
#UUID=5ee15463-49d1-403f-b77d-cb7038a909dc /myscripts ext4   defaults        0  2
#UUID=a51ca5c5-d563-4ad3-89a1-b715c377dfab /mintu     ext4   defaults        0  2
UUID="de33c6fb-692e-47fd-9f9f-6a629b8e291e" /mnt/raid-xfs xfs   defaults        0  0
[root@indlvrh91 ~]# mount -a
[root@indlvrh91 ~]#
```

---

## Step 6: Simulate Disk Failure and Rebuild (RAID 5)

### Mark one disk as failed:

```
#mdadm /dev/md0 --fail /dev/nvme0n3
```

This tells RAID:

"Pretend /dev/nvme0n3 has failed."

```
root@indlvrh91:~#
[root@indlvrh91 ~]# mdadm /dev/md0 --fail /dev/nvme0n3
mdadm: set /dev/nvme0n3 faulty in /dev/md0
[root@indlvrh91 ~]#
[root@indlvrh91 ~]#
```

### Verify the failure status:

```
#mdadm --detail /dev/md0
```

You'll see something like:

/dev/nvme0n3 -> faulty

/dev/nvme0n4 -> active

/dev/nvme0n5 -> active

and RAID will be in **degraded mode**.

```
[root@indlvrh91 ~]# [root@indlvrh91 ~]# mdadm --detail /dev/md0
/dev/md0:
      Version : 1.2
      Creation Time : Tue Oct 14 19:54:26 2025
      Raid Level : raid5
      Array Size : 1044480 (1020.00 MiB 1069.55 MB)
      Used Dev Size : 522240 (510.00 MiB 534.77 MB)
      Raid Devices : 3
      Total Devices : 3
      Persistence : Superblock is persistent

      Update Time : Wed Oct 15 11:08:26 2025
                  State : clean, degraded
      Active Devices : 2
      Working Devices : 2
      Failed Devices : 1
      Spare Devices : 0

      Layout : left-symmetric
      Chunk Size : 512K

Consistency Policy : resync

      Name : indlvrh91.na.xom.com:0  (local to host indlvrh91.na.xom.com)
      UUID : 02008a21:9cee95b9:5562918f:26cebbed
      Events : 20

      Number  Major  Minor  RaidDevice State
      -       0       0       0     removed
      1       259      6       1     active sync  /dev/nvme0n4
      3       259      7       2     active sync  /dev/nvme0n5

      0       259      5       -     faulty   /dev/nvme0n3
[root@indlvrh91 ~]# date
Wednesday 15 October 2025 11:09:39 AM IST
[root@indlvrh91 ~]# 
```

### Remove the failed device:

```
#mdadm /dev/md0 --remove /dev/nvme0n3
```

```
[root@indlvrh91 ~]# date
Wednesday 15 October 2025 11:09:39 AM IST
[root@indlvrh91 ~]# mdadm /dev/md0 --remove /dev/nvme0n3
mdadm: hot removed /dev/nvme0n3 from /dev/md0
[root@indlvrh91 ~]# mdadm --detail /dev/md0
/dev/md0:
      Version : 1.2
      Creation Time : Tue Oct 14 19:54:26 2025
      Raid Level : raid5
      Array Size : 1044480 (1020.00 MiB 1069.55 MB)
      Used Dev Size : 522240 (510.00 MiB 534.77 MB)
      Raid Devices : 3
      Total Devices : 2
      Persistence : Superblock is persistent

      Update Time : Wed Oct 15 11:10:31 2025
                  State : clean, degraded
      Active Devices : 2
      Working Devices : 2
      Failed Devices : 0
      Spare Devices : 0

      Layout : left-symmetric
      Chunk Size : 512K

Consistency Policy : resync

      Name : indlvrh91.na.xom.com:0  (local to host indlvrh91.na.xom.com)
      UUID : 02008a21:9cee95b9:5562918f:26cebbed
      Events : 21

      Number  Major  Minor  RaidDevice State
      -       0       0       0     removed
      1       259      6       1     active sync  /dev/nvme0n4
      3       259      7       2     active sync  /dev/nvme0n5
[root@indlvrh91 ~]# 
```

### (Optional) Re-add the same device to simulate disk replacement:

```
#mdadm /dev/md0 --add /dev/nvme0n3
```

This triggers a **rebuild** process.

```
[root@indlvrh91 ~]# mdadm /dev/md0 --add /dev/nvme0n3
mdadm: added /dev/nvme0n3
[root@indlvrh91 ~]# mdadm --detail /dev/md0
/dev/md0:
      Version : 1.2
      Creation Time : Tue Oct 14 19:54:26 2025
      Raid Level : raid5
      Array Size : 1044480 (1020.00 MiB 1069.55 MB)
      Used Dev Size : 522240 (510.00 MiB 534.77 MB)
      Raid Devices : 3
      Total Devices : 3
      Persistence : Superblock is persistent

      Update Time : Wed Oct 15 11:22:26 2025
      State : clean
      Active Devices : 3
      Working Devices : 3
      Failed Devices : 0
      Spare Devices : 0

      Layout : left-symmetric
      Chunk Size : 512K

Consistency Policy : resync

      Name : indlvrh91.na.xom.com:0  (local to host indlvrh91.na.xom.com)
      UUID : 02008a21:9cee95b9:5562918f:26cebbed
      Events : 84

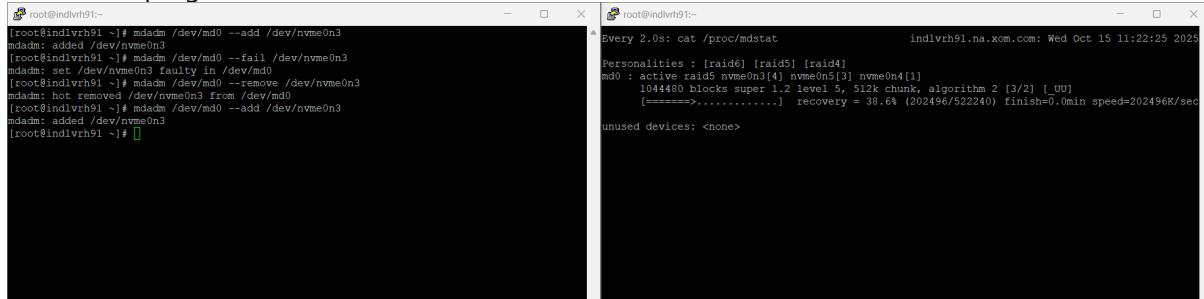
      Number  Major  Minor  RaidDevice State
         4      259      5          0     active sync   /dev/nvme0n3
         1      259      6          1     active sync   /dev/nvme0n4
         3      259      7          2     active sync   /dev/nvme0n5

[root@indlvrh91 ~]# date
Wednesday 15 October 2025 11:23:22 AM IST
[root@indlvrh91 ~]#
```

### Watch rebuild progress:

#watch cat /proc/mdstat

You'll see a progress bar as data is rebuilt.



### Verify everything is healthy again:

#mdadm --detail /dev/md0

#### Expected:

State : clean

Active Devices : 3

Working Devices : 3

Failed Devices : 0

```
[root@indlvrh91 ~]# mdadm --detail /dev/md0
/dev/md0:
      Version : 1.2
      Creation Time : Tue Oct 14 19:54:26 2025
      Raid Level : raid5
      Array Size : 1044480 (1020.00 MiB 1069.55 MB)
      Used Dev Size : 522240 (510.00 MiB 534.77 MB)
      Raid Devices : 3
      Total Devices : 3
      Persistence : Superblock is persistent

      Update Time : Wed Oct 15 11:22:26 2025
          State : clean
        Active Devices : 3
      Working Devices : 3
        Failed Devices : 0
        Spare Devices : 0

      Layout : left-symmetric
      Chunk Size : 512K

Consistency Policy : resync

      Name : indlvrh91.na.xom.com:0 (local to host indlvrh91.na.xom.com)
      UUID : 02008a21:9cee95b9:5562918f:26cebbed
      Events : 84

      Number  Major  Minor  RaidDevice State
        4      259      5        0     active sync   /dev/nvme0n3
        1      259      6        1     active sync   /dev/nvme0n4
        3      259      7        2     active sync   /dev/nvme0n5

[root@indlvrh91 ~]# date
Wednesday 15 October 2025 11:23:22 AM IST
[root@indlvrh91 ~]# 
```

### Important Note

| Rule  | Why   |
|---|---|
| Always fail/remove only <b>one disk</b> in RAID 5 | RAID 5 can survive <b>only one disk failure</b> |
| Don't reboot during rebuild                       | Rebuilding is intensive and must complete       |
| Use cat /proc/mdstat frequently                   | To monitor rebuild and ensure success           |

---

### Step 7: Performance Testing

#### 1. Write Test (Sequential Write Speed)

```
#cd /raid-xfs
#dd if=/dev/zero of=testfile bs=1M count=1000 oflag=direct
if=/dev/zero → reads zeroed bytes (simulates data write).
```

- of=testfile → file to write on RAID.
- bs=1M → block size = 1 MB.
- count=1000 → writes 1000 MB (1 GB).
- oflag=direct → bypass cache to measure actual disk speed.

After completion, it will show **MB/s write speed**.

```
[root@indlvrh91:~]# dd if=/dev/zero of=testfile bs=1M count=1000 oflag=direct
dd: error writing 'testfile': No space left on device
918+0 records in
917+0 records out
961544192 bytes (962 MB, 917 MiB) copied, 2.86243 s, 336 MB/s
[root@indlvrh91:~]# ls -lrth
total 917M
-rw-r--r--. 1 root root 917M Oct 14 20:26 testfile
[root@indlvrh91:~]# date
Tuesday 14 October 2025 08:26:40 PM IST
[root@indlvrh91:~]#
```

### Analysis of Output

```
dd: error writing 'testfile': No space left on device
918+0 records in
917+0 records out
961544192 bytes (962 MB, 917 MiB) copied, 2.93044 s, 328 MB/s
```

**No space left on device** → Expected, because your RAID 5 array is only ~1020 MB in total and you tried to write 1000 MB.

**Records in/out** → 917 full 1 MB blocks written

**Bytes copied** → 962 MB actually written.

**Speed** → ~328 MB/s sequential write speed.

This shows your RAID is functioning and the XFS filesystem is working properly.

## 2. Read Test (Sequential Read Speed)

```
#dd if=testfile of=/dev/null bs=1M iflag=direct
```

Reads the file from disk and discards it.

- iflag=direct → bypass cache to measure actual read speed.

```
[root@indlvrh91:~]#
[root@indlvrh91:~]# dd if=testfile of=/dev/null bs=1M iflag=direct
917+0 records in
917+0 records out
961544192 bytes (962 MB, 917 MiB) copied, 0.603097 s, 1.6 GB/s
[root@indlvrh91:~]# ls -lrth
total 917M
-rw-r--r--. 1 root root 917M Oct 14 20:26 testfile
[root@indlvrh91:~]# date
Tuesday 14 October 2025 08:26:57 PM IST
[root@indlvrh91:~]#
```

## 3. Clean up the test file (Clearly explained in step 8)

```
#rm -f /raid-xfs/testfile
```

Frees space after testing.

```
[root@indlvrh91 raid-xfs]# 
[root@indlvrh91 raid-xfs]# ls -lrth
total 917M
-rw-r--r--. 1 root root 917M Oct 14 20:26 testfile
[root@indlvrh91 raid-xfs]# rm -f /raid-xfs/testfile
[root@indlvrh91 raid-xfs]# ls -lrth
total 917M
-rw-r--r--. 1 root root 917M Oct 14 20:26 testfile
[root@indlvrh91 raid-xfs]# lsof | grep testfile
cat    6226          root   3r    REG      9,0 961544192      131 /mnt/raid-xfs/testfile
[root@indlvrh91 raid-xfs]# kill 6226
[root@indlvrh91 raid-xfs]# lsof | grep testfile
cat    6226          root   3r    REG      9,0 961544192      131 /mnt/raid-xfs/testfile
[root@indlvrh91 raid-xfs]# kill -9 6226
[root@indlvrh91 raid-xfs]# lsof | grep testfile
[1]+  Killed          cat testfile
[root@indlvrh91 raid-xfs]# rm -f /raid-xfs/testfile
[root@indlvrh91 raid-xfs]# ls -lrth
total 917M
-rw-r--r--. 1 root root 917M Oct 14 20:26 testfile
[root@indlvrh91 raid-xfs]# rm -f testfile
[root@indlvrh91 raid-xfs]# ls -lrth
total 0
[root@indlvrh91 raid-xfs]# df -h .
Filesystem      Size  Used Avail Use% Mounted on
/dev/md0       95G   39G  918M  5% /mnt/raid-xfs
[root@indlvrh91 raid-xfs]# date
Tuesday 14 October 2025 08:31:32 PM IST
```

#### 4. Use time for benchmarking

time sudo dd if=/dev/zero of=testfile bs=1M count=1000 oflag=direct  
Shows the total time it took to write, giving a clear throughput number.

```
[root@indlvrh91 ~]# time sudo dd if=/dev/zero of=testfile bs=1M count=1000 oflag=direct
1000+0 records in
1000+0 records out
1048576000 bytes (1.0 GB, 1000 MiB) copied, 15.3065 s, 68.5 MB/s
real    0m15.641s
user    0m0.071s
sys     0m0.343s
[root@indlvrh91 ~]# date
Wednesday 15 October 2025 11:33:18 AM IST
[root@indlvrh91 ~]# 
```

### Step 9: File Deletion & Process Lock Behavior

#### Scenario

Testing what happens when trying to delete a file that is still open by another process on an XFS filesystem mounted on RAID 5 (/mnt/raid-xfs).

#### Step-by-Step Analysis

##### 1. List directory contents

# ls -lrth

Shows one file named testfile, size **917 MB**.

Owner: root, Permissions: **rw-r--r--**

Confirms that the file exists in /mnt/raid-xfs/.

##### 2. Try to remove the file

#rm -f /raid-xfs/testfile

**Note:** you used /raid-xfs (missing /mnt/).

Since the mount point is /mnt/raid-xfs, this path may not actually exist, so the file remains untouched.

##### 3. Check open file handles

#lsof | grep testfile

#### Output:

cat 6226 root 3r REG 9,0 961544192 131 /mnt/raid-xfs/testfile

#### Meaning:

- **Process ID (PID):** 6226
- **Command:** cat
- **User:** root
- **File descriptor:** 3r (read-only)
- **Size:** ~917 MB

File is **currently open** by the cat process — this is why deletion doesn't work properly yet.

#### 4. Kill the process holding the file

```
#kill 6226
```

Terminates the cat process.

However, sometimes the process might linger briefly in memory, so you verify again.

#### 5. Recheck

```
#lsof | grep testfile
```

Still shows the same entry — meaning process is not yet fully gone.

#### 6. Force kill the process

```
#kill -9 6226
```

Sends a SIGKILL signal — **immediate termination**.

Process [1]+ Killed cat testfile confirms success.

#### 7. Reattempt file removal

```
#rm -f /mnt/raid-xfs/testfile
```

File successfully deleted this time.

Confirmed by ls -lrth → shows **total 0** (empty directory).

#### 8. Check filesystem usage

```
#df -h .
```

##### Output:

| Filesystem | Size | Used | Avail | Use% | Mounted on    |
|------------|------|------|-------|------|---------------|
| /dev/md0   | 956M | 9.5M | 947M  | 1%   | /mnt/raid-xfs |

Confirms that space used by the deleted file is now reclaimed.

The filesystem is healthy and mounted on **/mnt/raid-xfs**.

```
#date
```

Shows when this verification was done:

Tuesday 14 October 2025 08:31:32 PM IST

```
[root@indlvrh91 raid-xfs]#  
[root@indlvrh91 raid-xfs]# ls -lrth  
total 917M  
-rw-r--r--. 1 root root 917M Oct 14 20:26 testfile  
[root@indlvrh91 raid-xfs]# rm -f /raid-xfs/testfile  
[root@indlvrh91 raid-xfs]# ls -lrth  
total 917M  
-rw-r--r--. 1 root root 917M Oct 14 20:26 testfile  
[root@indlvrh91 raid-xfs]# lsof | grep testfile  
cat      6226          root    3r      REG            9,0  961544192      131 /mnt/raid-xfs/testfile  
[root@indlvrh91 raid-xfs]# kill 6226  
[root@indlvrh91 raid-xfs]# lsof | grep testfile  
cat      6226          root    3r      REG            9,0  961544192      131 /mnt/raid-xfs/testfile  
[root@indlvrh91 raid-xfs]# kill -9 6226  
[root@indlvrh91 raid-xfs]# lsof | grep testfile  
[1]+  Killed          cat testfile  
[root@indlvrh91 raid-xfs]# rm -f /raid-xfs/testfile  
[root@indlvrh91 raid-xfs]# ls -lrth  
total 917M  
-rw-r--r--. 1 root root 917M Oct 14 20:26 testfile  
[root@indlvrh91 raid-xfs]# rm -f testfile  
[root@indlvrh91 raid-xfs]# ls -lrth  
total 0  
[root@indlvrh91 raid-xfs]# df -h .  
Filesystem      Size   Used  Avail Use% Mounted on  
/dev/md0       956M   39M  918M   5% /mnt/raid-xfs  
[root@indlvrh91 raid-xfs]# date  
Tuesday 14 October 2025 08:31:32 PM IST
```

## Conclusion

We have set up a **RAID 5 + XFS setup on RHEL 9** is fully functional and stable.

- The **RAID array** has been successfully created, verified, and mounted.
- The **XFS filesystem** is operational, formatted, and tested for read/write performance.
- You've successfully demonstrated the **entire lifecycle** of a RAID setup:
  1. Disk verification and cleanup
  2. RAID creation with mdadm
  3. Filesystem creation (XFS)
  4. Mounting and verification
  5. Performance benchmarking
  6. Process/file lock handling
  7. Safe cleanup and teardown