

Complete Linux Server Monitoring & Troubleshooting Guide

*A Comprehensive Technical Reference for
System Administrators and DevOps Engineers*

Mohamed Achraf Sabbagh

December 5 , 2025

Topics Covered:

CPU Monitoring • Memory Management
Disk I/O Analysis • Network Diagnostics
Performance Optimization • Troubleshooting Workflows

This guide provides production-ready monitoring strategies,
real-world troubleshooting scenarios, and best practices
for maintaining high-performance Linux server environments.

Contents

1	Introduction	4
1.1	Purpose and Scope	4
1.2	Target Audience	4
1.3	How to Use This Guide	4
2	CPU Monitoring	5
2.1	Understanding CPU Architecture	5
2.2	Load Average Analysis	5
2.2.1	Command: uptime	5
2.2.2	Load Average Interpretation Matrix	5
2.3	Per-Core CPU Utilization	6
2.3.1	Command: mpstat	6
2.3.2	CPU Metrics Explained	6
2.4	Process-Level CPU Analysis	7
2.4.1	Interactive Monitoring with top	7
2.4.2	Sample Output Interpretation	7
2.5	CPU Saturation Detection	7
2.6	CPU Troubleshooting Workflow	8
3	Memory (RAM) Monitoring	9
3.1	Linux Memory Management Fundamentals	9
3.2	Primary Memory Analysis	9
3.2.1	Command: free -h	9
3.2.2	Column Interpretation	9
3.3	Memory Health Assessment	9
3.4	Identifying Memory-Hungry Processes	10
3.5	Swap Analysis	10
3.5.1	Understanding Swap Usage	10
3.6	Memory Troubleshooting Workflow	11
4	Disk & Storage Monitoring	12
4.1	Understanding Storage Performance	12
4.2	Disk Space Monitoring	12
4.2.1	Command: df -h	12
4.2.2	Filesystem Health Thresholds	12
4.3	Inode Usage Analysis	12
4.3.1	The Hidden Quota Problem	12
4.4	Disk I/O Performance	13
4.4.1	Command: iostat -x	13
4.4.2	Critical I/O Metrics	13
4.5	SMART Health Monitoring	14
4.5.1	Critical SMART Attributes	14
4.6	Disk Troubleshooting Workflow	15
5	Network Monitoring	16
5.1	Network Monitoring Layers	16
5.2	Interface Traffic Analysis	16
5.2.1	Command: ip -s link	16
5.2.2	Counter Interpretation	16

5.3	Connection Tracking	17
5.3.1	Command: ss (Socket Statistics)	17
5.3.2	TCP Connection States	17
5.4	Latency and Packet Loss Analysis	17
5.4.1	Basic Connectivity: ping	17
5.4.2	Advanced Path Analysis: mtr	18
5.5	Network Troubleshooting Workflow	18
6	Quick Reference Cards	20
6.1	CPU Quick Reference	20
6.2	Memory Quick Reference	20
6.3	Disk Quick Reference	21
6.4	Network Quick Reference	21
7	Conclusion	22

Abstract

This comprehensive guide provides system administrators, DevOps engineers, and infrastructure professionals with detailed methodologies for monitoring, diagnosing, and resolving performance issues in Linux server environments. Through practical examples, real-world scenarios, and systematic troubleshooting workflows, this document covers the four critical pillars of server health: CPU utilization, memory management, disk I/O performance, and network connectivity.

Each section includes command references, output interpretation, threshold guidelines, and actionable remediation strategies. The guide emphasizes holistic analysis, teaching readers to correlate metrics across subsystems to identify root causes rather than symptoms.

Keywords: Linux, System Administration, Performance Monitoring, Troubleshooting, DevOps, Infrastructure Management

1 Introduction

1.1 Purpose and Scope

In modern infrastructure environments, server performance directly impacts application availability, user experience, and business operations. This guide provides a systematic approach to:

- Monitor critical system resources in real-time
- Identify performance bottlenecks accurately
- Diagnose root causes of system degradation
- Implement effective remediation strategies
- Establish proactive monitoring frameworks

1.2 Target Audience

This guide is designed for:

- System Administrators managing Linux infrastructure
- DevOps Engineers responsible for application performance
- Site Reliability Engineers (SREs) maintaining production systems
- Infrastructure Engineers planning capacity
- Technical Support personnel diagnosing issues

1.3 How to Use This Guide

Key Information

Quick Reference: Each section includes:

- Command syntax and examples
- Output interpretation guidelines
- Health threshold matrices
- Real-world troubleshooting scenarios
- Remediation strategies

During an Incident: Jump directly to the relevant section's troubleshooting workflow.

For Learning: Read sequentially, practicing commands on test systems.

For Reference: Use the Quick Reference Cards at the end of each major section.

2 CPU Monitoring

2.1 Understanding CPU Architecture

Before implementing monitoring, understanding your system's CPU architecture is essential for accurate interpretation of metrics.

```

1 # Count total logical CPUs
2 nproc
3
4 # Detailed CPU information
5 lscpu
6
7 # Example output interpretation:
8 # CPU(s): 8           <- Total logical CPUs
9 # Thread(s) per core: 2 <- Hyperthreading enabled
10 # Core(s) per socket: 4 <- 4 physical cores
11 # Socket(s): 1        <- 1 physical CPU chip
12 # Result: 1 chip x 4 cores x 2 threads = 8 logical CPUs

```

Listing 1: CPU Architecture Discovery

Key Information

Why Architecture Matters:

A load average of 8.0 on an 8-core system indicates full utilization, while 16.0 indicates processes are queuing. Understanding this relationship is crucial for accurate capacity planning and performance assessment.

2.2 Load Average Analysis

Load average represents the average number of processes that are either running, runnable (waiting for CPU), or in uninterruptible sleep (waiting for I/O).

2.2.1 Command: uptime

```

1 $ uptime
2 14:32:18 up 23 days, 4:12, 3 users, load average: 2.15, 1.98, 1.76

```

Listing 2: Checking Load Average

The three numbers represent 1-minute, 5-minute, and 15-minute averages respectively.

2.2.2 Load Average Interpretation Matrix

Table 1: Load Average Assessment Guidelines

Load vs Cores	Meaning	Action Required
Load < Cores	Underutilized	Normal, capacity available
Load ≈ Cores	Fully utilized	Monitor trends
Load = 1.5× Cores	Moderate queuing	Investigate within minutes
Load = 2× Cores	Heavy queuing	Immediate investigation
Load > 3× Cores	Severe saturation	Emergency response

Example

Scenario A - 4-core system:

```
load average: 1.20, 1.45, 1.68
```

Analysis: All loads < 4.0 (healthy). Upward trend (1.20 → 1.68) indicates recent activity increase. Continue monitoring.

Scenario B - 4-core system:

```
load average: 12.50, 8.30, 4.20
```

Analysis:

- 1-min: 12.50 (3× cores) - **CRITICAL** spike occurring now
- 5-min: 8.30 (2× cores) - Heavy load past 5 minutes
- 15-min: 4.20 (1× cores) - Was normal 15 minutes ago

Action: Immediate investigation with `top` or `htop`.

2.3 Per-Core CPU Utilization

2.3.1 Command: mpstat

```
1 $ mpstat -P ALL 1 5
2 # -P ALL: show all processors
3 # 1: 1 second interval
4 # 5: 5 iterations
5
6 Average: CPU %usr %nice %sys %iowait %irq %soft %steal %idle
7 Average: all 15.23 0.00 4.12 2.34 0.00 0.50 0.00 77.81
8 Average: 0 45.67 0.00 6.21 0.45 0.00 0.12 0.00 47.55
9 Average: 1 12.34 0.00 3.88 1.22 0.00 0.45 0.00 82.11
10 Average: 2 10.11 0.00 3.45 3.56 0.00 0.67 0.00 82.21
11 Average: 3 9.87 0.00 3.02 4.11 0.00 0.56 0.00 82.44
```

Listing 3: Per-Core CPU Statistics

2.3.2 CPU Metrics Explained

%usr (User Space): Time executing user applications

- Normal: 20-70%
- High (>80%): CPU-intensive work (calculations, encoding)

%sys (System/Kernel): Time executing kernel operations

- Normal: 1-10%
- High (>20%): Excessive syscalls, context switching, or kernel work

%iowait: CPU idle waiting for I/O completion

- Normal: <5%
- Concerning: 10-20%

- Critical: >20% sustained

%steal (Virtual Machines Only): Time stolen by hypervisor

- Normal: <5%
- High (>10%): VM throttling or host oversold

Warning

Critical Indicator: If `%sys > %usr`, this often indicates a kernel bottleneck or inefficient I/O patterns. Investigate with `perf` or check connection counts and file operation patterns.

2.4 Process-Level CPU Analysis

2.4.1 Interactive Monitoring with top

```
1 top -o %CPU
2 # Press '1' to show individual cores
3 # Press 'H' to show threads
4 # Press 'c' to show full command line
5 # Press 'M' to sort by memory
```

Listing 4: Top Command Usage

2.4.2 Sample Output Interpretation

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
3421	www-data	20	0	982M	89M	12M	R	87.3	5.6	342:18	php-fpm: worker
5632	mysql	20	0	3.1G	1.8G	23M	S	23.5	12.1	1892:34	mysqld
7821	root	20	0	156M	45M	8M	R	15.2	2.8	45:23	python3 worker.py

Listing 5: Top Command Output

Key Columns:

- **%CPU:** Can exceed 100% (multi-threaded across cores)
- **TIME+:** Total CPU time consumed ($342:18 = 342$ minutes, 18 seconds)
- **S (State):** R=Running, S=Sleeping, D=Uninterruptible (I/O wait), Z=Zombie
- **RES:** Actual physical RAM used (most important memory metric)

2.5 CPU Saturation Detection

```
1 vmstat 1 5
2
3 procs -----memory----- --swap-- ---io-- -system- ---cpu---
4 r b swpd free buff cache si so bi bo in cs us sy id wa
5 3 2 524288 445632 125K 3.4G 0 0 12 456 890 1245 25 5 68 2
6 12 0 524288 398234 125K 3.4G 0 0 4 123 1123 1834 62 12 24 2
```

Listing 6: Checking Run Queue Depth

Critical Columns:

r (runnable): Processes waiting for CPU time

- Healthy: $r < \text{number of cores}$
- Saturated: $r > \text{cores}$ (Example: 12 runnable on 4-core = 8 queued)

b (blocked): Processes in uninterruptible sleep (I/O wait)

2.6 CPU Troubleshooting Workflow

1. **Quick Check:** `uptime` - Is load high?
2. **Identify Type:** `mpstat -P ALL 1 3` - CPU or I/O?
3. **Find Culprit:** `top -o %CPU` or `htop`
4. **Investigate Process:**

```
1 ps aux | grep <PID>          # Full command
2 lsof -p <PID>                # Open files/sockets
3 strace -c -p <PID>           # System call profile
4
```

5. **If I/O Wait High:** Proceed to Disk Monitoring section

3 Memory (RAM) Monitoring

3.1 Linux Memory Management Fundamentals

Key Information

Critical Concept: Linux aggressively uses free RAM for caching. A system showing "low free memory" is typically *healthy*, not problematic. The key metric is **available** memory, not **free**.

Memory Hierarchy:

1. Used by applications (cannot be reclaimed)
2. Cache/Buffers (kernel caching, can be freed instantly)
3. Free (truly unused)

3.2 Primary Memory Analysis

3.2.1 Command: free -h

```

1 $ free -h
2           total        used         free       shared    buff/cache   available
3 Mem:      15Gi       11Gi       500Mi      120Mi       3.5Gi       3.6
4           Gi
4 Swap:     2.0Gi      512Mi      1.5Gi

```

Listing 7: Memory Status Check

3.2.2 Column Interpretation

total: Physical RAM installed (15 GiB)

used: Memory used by applications + kernel (11 GiB)

free: Truly unused memory (500 MiB) - **IGNORE** for health

buff/cache: Kernel caching for performance (3.5 GiB) - **Good!**

available: **MOST CRITICAL** - Memory allocatable without swapping (3.6 GiB)

Swap used: Memory paged to disk (512 MiB)

3.3 Memory Health Assessment

Table 2: Memory Health Thresholds

Available %	Status	Action	Risk
> 20%	Healthy	Normal operation	None
10-20%	Monitor	Watch trends	Low
5-10%	Warning	Identify memory hogs	Medium
< 5%	Critical	Immediate action	High - OOM risk

Example

Scenario A - Healthy System:

```
Mem: 32Gi 28Gi 800Mi 200Mi 3.0Gi 3.5Gi (11% available)
```

Analysis: 28 GiB used, but 3.5 GiB available. Large cache (3.0 GiB) can be reclaimed.

Status: Healthy.

Scenario C - Memory Pressure:

```
Mem: 8Gi 7.8Gi 50Mi 100Mi 150Mi 180Mi (2% available)
Swap: 4Gi 2.1Gi 1.9Gi
```

Analysis: Only 180 MiB available, actively swapping 2.1 GiB. Status: **Critical** - Add RAM or reduce usage immediately.

3.4 Identifying Memory-Hungry Processes

```
1 $ ps aux --sort=-rss | head -n 10
2
3 USER      PID %CPU %MEM      VSZ      RSS COMMAND
4 mysql     1245  5.2 18.5 4523M 2987M /usr/sbin/mysqld
5 java      5632  3.1 15.2 8234M 2456M java -Xmx2G -jar app.jar
6 www       7821  1.2 12.3 2345M 1987M php-fpm: pool www
7 redis     3421  0.8  8.7 1234M 1405M redis-server *:6379
```

Listing 8: Top Memory Consumers

Key Metrics:

- **RSS (Resident Set Size):** Actual physical RAM used (most important)
- **VSZ (Virtual Size):** Usually irrelevant - includes shared libraries
- **%MEM:** Percentage of total RAM

3.5 Swap Analysis

3.5.1 Understanding Swap Usage

```
1 vmstat 1 5
2
3 procs -----memory----- --swap-- ---io---
4 r b    swpd   free   cache   si    so    bi    bo
5 1 0    524288 445632 3456K    0    0    12   456   <- No swapping (good)
6 1 0    524288 423156 3456K    0   234    8   234   <- Swapping out
7 1 2    524288 398234 3456K   456    0    4   123   <- Swapping in (BAD)
8 2 1    524288 375123 3456K   234   567   12   345   <- Both (CRITICAL)
```

Listing 9: Monitoring Swap Activity

Critical Columns:

si (swap in): KB/s read from swap to RAM - **MOST CRITICAL**

so (swap out): KB/s written from RAM to swap

Critical

Swap Thrashing: When `si` is high and sustained, the system spends more time swapping than working. This creates a "death spiral" where everything slows down exponentially. Users will report the system as "hanging" or "frozen."

3.6 Memory Troubleshooting Workflow

1. **Check Status:** `free -h` - Is available < 10%?
2. **Swap Activity:** `vmstat 1 5` - Is `si/so` > 0?
3. **Find Memory Hogs:** `ps aux -sort=-rss | head -n 20`
4. **Check for Leaks:**

```
1 watch -n 10 'ps aux | grep <process> | awk "{print \$6}"'
2 # If RSS continuously grows without plateau -> memory leak
3
```

5. **Review OOM History:** `dmesg | grep -i "killed process"`

4 Disk & Storage Monitoring

4.1 Understanding Storage Performance

Modern storage technologies have vastly different performance characteristics:

Table 3: Storage Performance Comparison

Type	IOPS	Sequential Throughput
HDD (7200 RPM)	80-160	100-200 MB/s
SATA SSD	10,000-90,000	500-600 MB/s
NVMe SSD	100,000-1,000,000+	3,000-7,000 MB/s

4.2 Disk Space Monitoring

4.2.1 Command: df -h

```

1 $ df -h
2 Filesystem      Size  Used  Avail Use% Mounted on
3 /dev/sda1        40G   28G   11G  73% /
4 /dev/sdb1        100G  92G   3.5G  96% /var
5 /dev/sdc1        500G  234G  241G  50% /home
6 tmpfs           7.8G  1.2G   6.6G  16% /run

```

Listing 10: Filesystem Space Usage

4.2.2 Filesystem Health Thresholds

Table 4: Disk Space Alert Levels

Use%	Status	Action	Risk
< 70%	Healthy	Monitor trends	None
70-85%	Elevated	Plan cleanup/expansion	Low
85-95%	Warning	Act within days	Service failures
95-98%	Critical	Immediate action	High failure risk
> 98%	Emergency	Urgent intervention	System instability

Warning

Why Filesystems Fail Before 100%:

- Reserved blocks (5% for root on ext4 default)
- Metadata overhead (inodes, journals, allocation tables)
- Application behavior (many apps fail ungracefully when writes fail)

4.3 Inode Usage Analysis

4.3.1 The Hidden Quota Problem

```

1 $ df -i
2 Filesystem      Inodes   IUsed   IFree  IUse% Mounted on
3 /dev/sda1        6553600 340000 6213600     6% /
4 /dev/sdb1        6553600 6553590        10 100% /var
5 /dev/sdc1        32768000 123456 32644544     1% /home

```

Listing 11: Checking Inode Usage

Critical**The 100% Inode Problem:****Filesystem:** /dev/sdb1 has 55GB free space but 100% inodes used.**Result:** Cannot create ANY new files despite available space!**Error Message:**

```
$ touch test.txt
touch: cannot touch 'test.txt': No space left on device
```

Common Causes:

- Mail queues with millions of small files
- Many rotated log files (.gz archives)
- PHP/web session directories
- Node.js/Python cache directories

4.4 Disk I/O Performance

4.4.1 Command: iostat -x

```

1 $ iostat -x 1 5
2
3 Device: r/s   w/s   rkB/s  wkB/s  avgrrq-sz  avgqu-sz  await  svctm %util
4 sda    12.3  45.6   1234   8765      234.5       2.34   45.2    8.9   67.5
5 sdb   456.7 23.4  45678   2345      987.6      12.45  234.6    2.1   98.9
6 sdc     0.5   1.2     50    102      128.0       0.12   3.2    0.5    5.2

```

Listing 12: Detailed I/O Statistics

4.4.2 Critical I/O Metrics

r/s, w/s: Reads/writes per second (IOPS)**avgqu-sz:** Average queue length - Key saturation metric

- < 1: Disk keeping up
- 1-4: Moderate queue, acceptable
- > 4: Requests backing up
- > 10: Severe congestion

await: Average time (ms) for I/O requests

- HDD baseline: 5-15ms

- SSD baseline: <1ms
- > 20ms on SSD: Problem
- > 50ms on HDD: Severe problem
- > 100ms: Critical - users notice slowness

%util: Percentage of time with at least one I/O pending

- 100% on SSD ≠ saturated (handles parallel I/O)
- 100% on HDD = likely saturated

Example

Example 1: HDD Saturation

```
Device: r/s    w/s    avgqu-sz   await   %util
sda      234.5   89.3     8.45    156.7   98.5
```

Analysis:

- avgqu-sz=8.45: Many requests waiting
- await=156.7ms: Very slow (normal HDD 10ms)
- %util=98.5%: Saturated

Diagnosis: HDD cannot handle 323 IOPS (r/s + w/s)

Solutions: Migrate to SSD, optimize queries, RAID configuration, move I/O-heavy workloads

4.5 SMART Health Monitoring

```
1 # Quick health check
2 smartctl -H /dev/sda
3
4 # Full SMART data
5 smartctl -a /dev/sda
```

Listing 13: Checking Disk Health

4.5.1 Critical SMART Attributes

Table 5: Critical SMART Indicators

Attribute	Threshold	Action
Reallocated_Sector_Ct	> 0	Plan replacement
	> 10	Imminent failure
Current_Pending_Sector	> 0	Monitor closely
	> 5	Replace soon
Offline_Uncorrectable	> 0	Data loss risk
Reported_Uncorrect	> 0	Reliability concern
Power_On_Hours	> 40,000	Consider age in planning

Critical**Immediate Replacement Indicators:**

- Reallocated_Sector_Ct > 10
- Current_Pending_Sector > 5
- Offline_Uncorrectable > 0
- Any combination of growing errors

Action: Immediate backup, snapshot data, and schedule replacement.

4.6 Disk Troubleshooting Workflow

1. **Check Space:** df -h and df -i
2. **If Space Issue:**

```
1 du -sh /* | sort -rh
2 find / -type f -size +1G -exec ls -lh {} \;
3
```
3. **Check I/O:** iostat -x 1 5
4. **If High I/O:** iotop -o -a
5. **Check Health:** smartctl -a /dev/sda

5 Network Monitoring

5.1 Network Monitoring Layers

Network troubleshooting requires a layered approach:

1. **Physical:** Cables, NICs, link speed
2. **Data Link:** Ethernet frames, errors, drops
3. **Network:** IP packets, routing, latency
4. **Transport:** TCP connections, UDP packets
5. **Application:** HTTP, SSH, Database protocols

5.2 Interface Traffic Analysis

5.2.1 Command: ip -s link

```

1 $ ip -s link show eth0
2 2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500
3   link/ether 00:1a:2b:3c:4d:5e brd ff:ff:ff:ff:ff:ff
4     RX: bytes packets errors dropped overrun mcast
5       98765432109 87654321 0      523      0      12345
6     TX: bytes packets errors dropped carrier collsns
7       45678901234 56789012 0      127      0      0

```

Listing 14: Interface Statistics

5.2.2 Counter Interpretation

errors: Malformed frames (CRC, alignment, framing)

- $> 0.01\%$ of packets: Investigate physical layer
- Causes: Bad cable, faulty NIC, switch port issue

dropped: Packets discarded by kernel

- RX dropped: Receive buffer overflow
- TX dropped: Transmit queue full
- > 0 : Tune ring buffers or investigate driver

overrun: NIC buffer overflow

- > 0 : Severe hardware bottleneck (rare)

carrier: Carrier signal lost (Layer 1 issue)

- > 0 : Check physical connection, cable, transceiver

5.3 Connection Tracking

5.3.1 Command: ss (Socket Statistics)

```

1 $ ss -s
2 Total: 1247 (kernel 1532)
3 TCP:    456 (estab 387, closed 45, orphaned 2, timewait 38)
4 UDP:    89
5
6 # List listening ports
7 $ ss -tlnp
8 State  Recv-Q Send-Q Local Address:Port      Process
9 LISTEN  0        128    0.0.0.0:22          users:(("sshd",pid=1234))
10 LISTEN  0        128    0.0.0.0:80         users:(("nginx",pid=5678))
11 LISTEN  0        80     127.0.0.1:3306       users:(("mysqld",pid=9012))

```

Listing 15: Connection Summary

5.3.2 TCP Connection States

LISTEN: Server socket waiting for connections

- 0.0.0.0 = listening on all interfaces
- 127.0.0.1 = localhost only (secure)

ESTAB (Established): Active connection

- High count may indicate heavy traffic or connection leak

TIME-WAIT: Connection closed, waiting for final ACK

- Normal: lasts 60 seconds ($2 \times \text{MSL}$)
- Many TIME-WAIT (thousands): High connection churn
- Solution: Tune `net.ipv4.tcp_tw_reuse=1`

CLOSE-WAIT: Remote closed, local app hasn't closed

- Growing CLOSE-WAIT: Application bug (not closing sockets)
- Solution: Fix application code, restart service

5.4 Latency and Packet Loss Analysis

5.4.1 Basic Connectivity: ping

```

1 $ ping -c 10 -i 0.2 8.8.8.8
2 PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
3 64 bytes from 8.8.8.8: icmp_seq=1 ttl=118 time=12.2 ms
4 64 bytes from 8.8.8.8: icmp_seq=2 ttl=118 time=11.8 ms
5 ...
6 --- 8.8.8.8 ping statistics ---
7 10 packets transmitted, 10 received, 0% packet loss
8 rtt min/avg/max/mdev = 11.756/12.437/13.145/0.421 ms

```

Listing 16: Latency Testing

Latency Guidelines:

- Local network: <1ms
- Same city: 1-10ms
- Same country: 10-50ms
- Intercontinental: 100-300ms

Packet Loss Impact:

- 0%: Perfect
- <1%: Acceptable
- 1-5%: Noticeable degradation (VoIP suffers)
- >5%: Severe problems

5.4.2 Advanced Path Analysis: mtr

```

1 $ mtr -r -c 100 -n example.com
2
3 HOST           Loss% Snt  Last   Avg   Best Wrst StDev
4 1. 10.0.0.1    0.0% 100   0.4   0.5   0.3  1.2  0.1
5 2. 203.0.113.1 0.0% 100  10.2  10.5  9.8 15.3  0.8
6 3. 198.51.100.45 2.0% 100  20.1  22.3 19.5 80.4  8.7
7 4. 192.0.2.12   2.0% 100  21.5  23.1 20.2 82.1  9.1
8 5. 93.184.216.34 0.0% 100  22.3  23.5 21.8 28.9  1.2

```

Listing 17: Multi-hop Trace Route

Key Information

Interpreting mtr Results:

Loss at single hop, but not later: Often false positive (router rate-limiting ICMP)

Loss at hop AND all subsequent hops: Real problem at that hop

In example above: 2% loss starts at hop 3 and persists → Root cause at hop 3

5.5 Network Troubleshooting Workflow

1. Physical Layer:

```

1 ethtool eth0          # Link up? Speed correct?
2 ip link show eth0     # Interface UP?
3

```

2. Data Link Layer:

```

1 ip -s link show eth0      # Errors? Drops?
2 ethtool -S eth0          # Detailed NIC stats
3

```

3. Network Layer:

```

1 ip route                # Default route correct?
2 ping 8.8.8.8             # Internet reachable?
3 ping <gateway>          # Local gateway reachable?
4

```

4. Transport Layer:

```
1 ss -s                      # Connection states normal?  
2 ss -tlnp                   # Services listening?  
3
```

5. Application Layer:

```
1 curl -I http://localhost    # App responding?  
2 journalctl -u <service>    # App logs  
3
```

6 Quick Reference Cards

6.1 CPU Quick Reference

Commands:

```
uptime          # Load averages
top             # Interactive monitor
htop            # Better alternative
mpstat -P ALL 1 # Per-core stats
ps aux --sort=-%cpu # Top CPU processes
```

Key Metrics:

- Load < cores = OK
- %iowait > 10% = check disk
- %steal > 10% = VM throttled

6.2 Memory Quick Reference

Commands:

```
free -h          # Memory overview
ps aux --sort=-rss # Top memory users
vmstat 1 5       # Swap activity
smem -tk         # Detailed breakdown
```

Key Metrics:

- available < 10% = warning
- swap si/so > 0 = problem
- OOM kills = critical

6.3 Disk Quick Reference

Commands:

```
df -h          # Disk space
df -i          # Inode usage
du -sh /*      # Directory sizes
iostat -x 1 5  # I/O performance
iostop -o       # Top I/O processes
smartctl -a /dev/sda  # SMART health
```

Key Metrics:

- Space > 90% = act soon
- await > 20ms SSD = problem
- await > 50ms HDD = problem
- %util 100% HDD = saturated

6.4 Network Quick Reference

Commands:

```
ss -s          # Connection summary
ss -tlnp        # Listening services
ip -s link      # Interface stats
iftop -i eth0    # Bandwidth per connection
mtr <destination>  # Path analysis
ping -c 10 <host>  # Basic connectivity
```

Key Metrics:

- Packet loss > 1% = issue
- Latency variation = jitter
- errors/drops > 0 = investigate
- TIME-WAIT high = tune reuse

7 Conclusion

Effective Linux server monitoring and troubleshooting requires:

- **Systematic methodology** - Follow structured workflows rather than random investigation
- **Holistic analysis** - Correlate metrics across CPU, memory, disk, and network
- **Understanding context** - Know your baseline and recognize deviations
- **Proactive monitoring** - Detect issues before they impact users
- **Continuous learning** - Document incidents and build organizational knowledge

*Remember: The goal is not just to fix problems,
but to understand systems deeply enough
to prevent problems before they occur.*

Appendix A: Essential Commands Cheat Sheet

Table 6: *
CPU Monitoring Commands

Command	Purpose
<code>uptime</code>	Display load averages
<code>top</code>	Interactive process viewer
<code>htop</code>	Enhanced interactive viewer
<code>mpstat -P ALL 1</code>	Per-core CPU statistics
<code>ps aux -sort=-%cpu</code>	List processes by CPU usage
<code>vmstat 1 5</code>	System statistics including run queue
<code>pidstat 1</code>	Per-process statistics over time
<code>perf top</code>	Real-time profiling

Table 7: *
Memory Monitoring Commands

Command	Purpose
<code>free -h</code>	Memory usage summary
<code>vmstat 1 5</code>	Memory and swap statistics
<code>ps aux -sort=-rss</code>	List processes by memory usage
<code>pmap -x <PID></code>	Memory map of specific process
<code>smem -tk</code>	Detailed memory breakdown
<code>slabtop</code>	Kernel slab cache information
<code>cat /proc/meminfo</code>	Detailed memory statistics

Table 8: *
Disk Monitoring Commands

Command	Purpose
<code>df -h</code>	Filesystem space usage
<code>df -i</code>	Inode usage
<code>du -sh /*</code>	Directory sizes
<code>iostat -x 1 5</code>	Disk I/O statistics
<code>iotop -o</code>	Per-process I/O usage
<code>smartctl -a /dev/sda</code>	SMART disk health data
<code>lsblk</code>	List block devices
<code>blkid</code>	Block device attributes

Table 9: *
Network Monitoring Commands

Command	Purpose
<code>ss -s</code>	Socket statistics summary
<code>ss -tlnp</code>	Listening TCP sockets
<code>ip -s link</code>	Interface statistics
<code>ip route</code>	Routing table
<code>ethtool eth0</code>	NIC settings and status
<code>iftop -i eth0</code>	Real-time bandwidth usage
<code>nethogs</code>	Per-process bandwidth
<code>mtr <host></code>	Network path analysis
<code>tcpdump -i eth0</code>	Packet capture
<code>nmap</code>	Network scanning